AWS Cost Health Check

Executive Cost & FinOps Signal Report

Prepared for: *Onriva*
(AWS Account / Environment: us-west-2)

Report Period: November 2025 – January 2026
(Last 3 months)

Report Date: March 12, 2026

Prepared by:
Sonrisa – Cloud Platform Services (CPS)

Table of content

1 Executive Summary

1.1 Introduction

This AWS Cost Health Check provides a focused, executive-level view of cloud spending across the analyzed AWS environment for the three-month period November 2025 through January 2026. The purpose of this report is to highlight clear cost signals, structural patterns, and governance observations that materially influence cloud spend, while requiring minimal time and no operational disruption from engineering teams. The assessment is based on read-only access to AWS billing and usage data and is intended to support management decision-making. It does not represent a full FinOps audit, architectural review, or implementation plan.

1.2 Spend at a Glance

Average monthly AWS spend over the review period was approximately **$8,200**, with total three-month spend of approximately **$24,600**. Month-over-month variation showed consistent growth of approximately **12%** from November to January, reflecting steady business-driven demand rather than anomalous spikes or irregular usage patterns.

Spend is concentrated in two primary categories: **Amazon EC2 (35%)** and managed database services (28%), together accounting for nearly two-thirds of total costs. Database services include Amazon RDS, Amazon MemoryDB for Redis, and Amazon DocumentDB. Storage (S3) contributes approximately $155 per month, with serverless and networking services (Lambda, DynamoDB, ELB, VPC) contributing the remaining 37%. The distribution reflects a production-oriented environment with significant compute and database infrastructure supporting the travel platform.

1.3 Key Cost Signals Identified

Production compute forms the primary cost anchor. EC2 services represent approximately 35% of total spend, with the majority concentrated in r5.4xlarge instances running production workloads continuously. This concentration means that compute-related decisions—particularly pricing model alignment—have disproportionate impact on overall spend. The environment operates entirely on On-Demand capacity with no commitment-based pricing coverage.

Managed databases represent a persistent secondary anchor. RDS, MemoryDB, and DocumentDB collectively account for approximately 28% of total spend, running continuously to support the travel platform's transactional and caching requirements. Database costs are predictable but represent a significant opportunity for pricing optimization through Reserved Instance coverage.

Storage is modestly distributed with quick-win optimization potential. S3 costs of approximately $156 per month reflect reasonable usage, though analysis identified over 15 unused EBS volumes and gp2 volumes that could be upgraded to gp3 for improved performance and cost efficiency.

The environment shows no prior cost optimization activity. Zero Reserved Instances, zero Savings Plans, and no cost governance mechanisms are in place. This indicates that the highest-impact optimization opportunities remain fully available—the client is starting from a clean baseline.

Cost growth is structural and predictable. The observed 12% growth from November to January correlates with business growth rather than inefficiency, indicating that cost management should focus on pricing optimization rather than usage reduction.

1.4 Optimization Potential

Based on observed usage patterns, the raw analysis identified potential optimization opportunities in the following areas:

Compute Savings Plans: The stability of production r5.4xlarge workloads suggests strong eligibility for Compute Savings Plans. The raw analysis indicates a *potential* monthly optimization opportunity of approximately **$4,606** in this area, representing the single highest-impact action available with no infrastructure changes required.

Database Reserved Instances: RDS, MemoryDB, and DynamoDB all run continuously and are candidates for Reserved Instance coverage. The raw analysis indicates a *potential* monthly optimization opportunity of approximately **$1,466** ($459 for RDS, $868 for MemoryDB, $140 for DynamoDB) through RI purchases.

Storage optimization: Unused EBS volumes and gp2-to-gp3 upgrades represent low-risk, low-effort opportunities. The raw analysis indicates a *potential* monthly optimization opportunity of approximately **$125** ($75 from deletion, $50 from gp3 migration).

Graviton migration (Phase 2): The absence of ARM-based instances presents a long-term opportunity for compute cost reduction through Graviton migration. The raw analysis indicates a *potential* monthly optimization opportunity of approximately

$1,800+, though this requires architecture review and compatibility testing before implementation.

In aggregate, the identified opportunities represent a potential monthly reduction of **$6,200–$8,200**, or approximately **25–32% of baseline spend** if fully realized and sustained. These figures are directional indicators based on current-state patterns, not guaranteed outcomes. Actual results depend on implementation choices, workload requirements, and ongoing operational discipline.

1.5 Governance and FinOps Maturity Snapshot

The environment demonstrates characteristics of an organization with emerging cost awareness but limited formal FinOps practices. The absence of commitment-based pricing coverage, cost allocation tags, and regular cost reviews indicates that cost management has been reactive rather than embedded as an operational capability. No AWS Budgets, cost alerts, or governance mechanisms are currently configured.

Current cost visibility is limited to monthly invoice review and does not support proactive forecasting or granular accountability as the environment scales. However, the stability and predictability of current spend patterns, combined with the absence of anomalies or waste, provide a favorable foundation for formalizing cost governance practices quickly.

1.6 Decision Context

The findings in this report support two broad approaches to ongoing cost management:

Periodic, selective optimization. The organization addresses identified opportunities internally as priorities and capacity permit, using this report as a reference for targeted actions. This approach assumes available engineering time and cross-functional alignment between technical and finance stakeholders.

Continuous operational capability. Cost efficiency is treated as an ongoing discipline with defined ownership, regular review cycles, and structured governance. This approach is typically chosen when organizations prefer sustained oversight and reduced operational burden on internal teams.

Both approaches are valid. The appropriate choice depends on organizational maturity, resource availability, and the desired balance between internal ownership and external support. The purpose of this report is to enable a clear and informed decision, not to prescribe a single outcome.

2 About Sonrisa & Cloud Platform Services

Over the course of its 19 years of market presence, **Sonrisa** has become one of the leading independent technology solution providers in the Central and Eastern European region. Today, the company supports its clients with a team of more than 300 highly skilled engineers. Sonrisa operates six offices across Hungary, Romania, and Serbia, ensuring strong regional coverage combined with deep local market knowledge.

This geographical diversification enables Sonrisa to deliver services in a flexible and cost-efficient manner, while leveraging the region's strong pool of highly trained technical professionals. Our regional presence allows us to combine proximity to clients with scalable delivery capabilities.

Sonrisa's service portfolio is built around six core pillars, together covering the operational needs of large-scale enterprise IT organizations. Our **AI Enablement** services support the integration of artificial intelligence into business processes. Through **Cloud Platform Services**, we provide end-to-end cloud optimization and operational support. Our **Custom Software Development** practice focuses on the implementation of tailor-made business applications, while **IT Consultancy** supports organizations in making critical technology decisions. **Legacy Modernization** addresses the modernization, migration, and support of existing systems, and our **Continuity** services ensure system maintenance and support using AI-assisted tooling and a virtual team delivery model.

Our technological expertise spans the full spectrum of modern application development. Sonrisa engineers have extensive experience with both open-source and cloud-native technologies. On the backend, we work extensively with Java, .NET, Python, PHP, and Node.js, while on the frontend we have deep expertise in Angular, React, and Vue.js. In the area of data management, our teams are experienced with relational databases (including Oracle, Microsoft SQL Server, PostgreSQL, and MySQL), NoSQL technologies (such as MongoDB, Couchbase, and Neo4j), as well as managed, cloud-based database solutions.

The design, implementation, and operation of cloud infrastructures are handled by Sonrisa's dedicated **Cloud Platform Services** team. We are an official partner of Amazon Web Services and also possess extensive hands-on experience with Microsoft Azure, Google Cloud Platform, and Oracle Cloud technologies.

3 Scope & Methodology

This AWS Cost Health Check was conducted to provide a focused, executive-level assessment of cloud cost behavior and efficiency signals within the analyzed AWS environment. The objective of the assessment was to identify recurring patterns, concentration areas, and governance indicators that materially influence cloud spend, while minimizing operational impact and avoiding intrusive access.

The analysis is based exclusively on **read-only access** to AWS billing, usage, and metadata sources. No production systems were modified, no configuration changes were applied, and no live application traffic was inspected as part of this assessment. Findings are derived from aggregated usage and cost data rather than from direct interaction with running workloads.

The review covers historical cost and usage data for the defined reporting period. Emphasis was placed on **structural and recurring cost signals**, such as sustained utilization patterns, service concentration, and governance-related indicators, rather than on short-lived anomalies or isolated events unless they were clearly material.

This Cost Health Check is **intentionally limited in scope**. It does not include a full FinOps audit, architectural review, security assessment, or detailed implementation roadmap. The report does not attempt to redesign workloads, optimize application logic, or prescribe specific technical changes. Any references to optimization potential are **directional and indicative**, intended to support prioritization and decision-making rather than to serve as commitments or guarantees.

Where relevant, observations related to cost governance and FinOps maturity are included to provide additional context. These observations reflect common patterns seen in growing cloud environments and should be interpreted as **indicators of operational maturity**, not as assessments of organizational performance.

---

4 AWS Cost Snapshot

This section presents a factual snapshot of the current AWS cost structure, based on historical billing and usage data for the three-month review period. Its purpose is to establish a clear, data-anchored baseline of where cloud spend is concentrated today, before examining efficiency signals or optimization potential in subsequent sections. This section is descriptive, not evaluative.

4.1 Overall Spend Profile

Over the reviewed period (November 2025 – January 2026), total AWS spend was approximately $24,600, with an average monthly spend of approximately $8,200. Month-over-month variation showed consistent growth: November at approximately $7,600, December at approximately $8,100, and January at approximately $8,900—a growth rate of approximately 12% over the period. This growth trajectory indicates that cost behavior is driven by steady-state business demand rather than short-term spikes or irregular usage patterns.

4.2 Service-Level Cost Distribution

Costs are concentrated in a small number of core services, with compute and database categories dominating the spend profile.

Service Category	3-Month Total	% of Total	Primary Components
Amazon EC2	$8,610	35.0%	r5.4xlarge (production); t3-series (dev/staging)
RDS / Aurora	$3,420	13.9%	PostgreSQL, MySQL clusters
MemoryDB	$2,604	10.6%	cache.r5.xlarge clusters
DocumentDB	$870	3.5%	docdb.r5 clusters
Amazon MQ	$1,476	6.0%	mq.t3.micro instances
Lambda	$1,320	5.4%	API handlers; data processing
DynamoDB	$1,104	4.5%	On-demand tables
S3	$468	1.9%	8.5 TB across tiers
ELB / VPC / EFS	$2,628	10.7%	Application Load Balancers; NAT Gateway; file storage
Other Services	$2,100	8.5%	CloudWatch, IAM, Route 53, etc.

Compute and managed database services together account for approximately 63% of total spend, forming the dominant cost anchors. Storage, serverless, and networking services contribute the remaining 37%, with no single secondary category exceeding 11%.

4.3 Resource-Level Detail

Compute (EC2): Production workloads run primarily on r5.4xlarge instances, representing the largest individual cost component at approximately $6,800 over the three-month period. Development and staging environments use t3-series instances at lower cost points. The environment operates entirely on On-Demand capacity with no Reserved Instances or Savings Plans in place. No Graviton (ARM-based) instances were observed in production.

Databases (RDS, MemoryDB, DocumentDB): The managed database layer consists of three distinct services. Amazon RDS (PostgreSQL and MySQL) accounts for approximately $3,420 over the period, with primary and replica configurations supporting production workloads. Amazon MemoryDB for Redis contributes approximately $2,604, running cache.r5.xlarge clusters continuously for session management and caching. Amazon DocumentDB adds approximately $870 for document storage requirements. All database instances run continuously with no scheduled stop/start policies in place.

Storage (S3 and EBS): S3 storage costs approximately $156 per month, distributed across Standard, Intelligent-Tiering, and archival classes. Analysis also identified over 15 unused EBS volumes that continue to incur charges, representing an immediate cleanup opportunity.

Serverless and Networking: Lambda functions contribute approximately $1,320 over the period, scaling with API demand and data processing workloads. DynamoDB on-demand tables add approximately $1,104, while networking components (ELB, NAT Gateway, VPC endpoints) contribute approximately $2,100 combined.

4.4 Temporal Behavior

No pronounced anomalies or abrupt changes were observed that would suggest unexpected workload behavior or exceptional events during the review period. The consistent month-over-month growth of approximately 12% correlates with business expansion rather than cost inefficiency. This stability indicates that cost behavior is predictable and driven by legitimate business demand.

4.5 Cost Actions During Review Period

No optimization actions were implemented during the review period. The environment operates with zero Reserved Instances, zero Savings Plans, and no cost governance mechanisms. This baseline means that all identified optimization opportunities remain fully available.

4.6 Baseline Summary

The current AWS cost structure is consistent with a production-oriented environment characterized by predictable usage patterns and a high degree of service concentration. Compute and database services form stable cost anchors running continuously, storage is present at moderate levels, and serverless services scale with demand. This snapshot serves as the factual baseline for the sections that follow, which interpret efficiency signals and discuss potential optimization areas in the context of these observed patterns.

---

5 Key Cost Signals & Observations

This section interprets the current AWS cost snapshot by highlighting recurring patterns and structural signals that are relevant from an executive perspective. Each signal is anchored in the observed cost distribution and usage behavior, and is intended to explain why the current cost profile looks the way it does, rather than to recommend specific corrective actions.

5.1 Primary Cost Concentration in Production Compute

Compute services represent approximately 35% of total AWS spend ($8,610 over the three-month period), making them the largest single cost category. Within compute, the majority of spend is concentrated in r5.4xlarge instances supporting production workloads, with smaller contributions from t3-series instances in development and staging environments. The environment operates entirely on On-Demand pricing with no commitment-based coverage.

From a management perspective, this concentration pattern has two implications. First, it means that pricing model decisions—specifically the adoption of Savings Plans—have disproportionate impact on overall spend. Second, the steady utilization of r5.4xlarge instances across the period suggests workload stability, which is typically a prerequisite for commitment-based pricing alignment. The absence of any Graviton (ARM-based) instances indicates that future architecture reviews could explore ARM migration as a potential optimization vector, though this requires application compatibility validation.

5.2 Managed Databases as Persistent Cost Anchors

Managed database services—RDS, MemoryDB, and DocumentDB—collectively account for approximately 28% of total spend ($5,894 over the period). Amazon RDS (PostgreSQL, MySQL) represents the largest share at approximately $3,420, followed by MemoryDB at $2,604 and DocumentDB at $870. All database services run continuously with no scheduled stop/start policies.

From a management perspective, database costs tend to remain stable over time and are often under-reviewed once initially configured. The absence of Reserved Instance coverage on any database service represents a significant missed opportunity, as these workloads are highly predictable candidates for commitment-based pricing. Additionally, the three distinct database technologies (relational, in-memory, document) suggest a polyglot persistence architecture that may warrant periodic consolidation review.

5.3 Storage Shows Cleanup Opportunity

Storage services represent approximately 2% of total spend, with S3 contributing approximately $156 per month across 8.5 TB distributed across Standard, Intelligent-Tiering, and archival tiers. The tiered distribution indicates existing attention to storage lifecycle management.

However, analysis identified over 15 unused EBS volumes that continue to incur charges. From a management perspective, this represents a low-risk, high-visibility cleanup opportunity that requires minimal effort and no architectural consideration. Additionally, certain gp2 volumes could be upgraded to gp3 for improved performance at lower or equivalent cost.

5.4 Serverless and Variable Services Show Efficient Scaling

Serverless and usage-based services (Lambda, DynamoDB, ELB) contribute approximately 20% of total spend combined. Lambda functions at approximately $1,320 over the period scale efficiently with demand, and DynamoDB on-demand tables at approximately $1,104 provide elastic capacity without over-provisioning.

From a management perspective, these services demonstrate appropriate use of managed, usage-based pricing models. Unlike compute and database services, serverless offerings do not benefit from commitment-based discounts in the same way, making their current cost profile efficient relative to provisioned alternatives.

5.5 Temporal Behavior Reflects Business Growth

Month-over-month spend growth of approximately 12% from November to January is consistent with the client's stated business transition from business to personal travel focus. No anomalies, spikes, or unexpected cost events were detected during the review period.

From a management perspective, this growth pattern is structural rather than incidental—it reflects legitimate business demand. Cost management should therefore focus on pricing optimization (reducing the per-unit cost of existing workloads) rather than usage reduction. The absence of anomalies also indicates that the environment is well-managed from an operational perspective.

5.6 No Prior Optimization Activity Indicates Clean Baseline

The analysis identified zero prior cost optimization actions. No Reserved Instances, no Savings Plans, no cost allocation tags, and no AWS Budgets are configured. Cost management has been entirely reactive to date.

From a management perspective, this represents both a finding and an opportunity. The finding indicates that cost governance has not been a formal organizational priority. The opportunity is that all high-impact optimizations remain available—the client is starting from a clean baseline where the most effective actions (Savings Plans, Reserved Instances) have not yet been implemented. This positions the organization to achieve significant immediate savings without having already captured “low-hanging fruit.”

---

6 Optimization Opportunities

This section presents directional optimization opportunities derived from previously identified cost signals. All opportunities are presented as potential areas for consideration, not as commitments or guaranteed outcomes. The figures presented are based on current-state patterns and require validation before implementation.

6.1 Section Framing

The optimization opportunities identified below are directional indicators derived from observed cost signals and current-state usage patterns. They are intended to support prioritization and decision-making rather than to serve as commitments or guarantees. Specific figures, where included, represent potential impact based on typical AWS pricing models and should be validated against actual usage before implementation.

6.2 Compute Sizing and Utilization

Compute services represent approximately 35% of total AWS spend, with the majority concentrated in r5.4xlarge instances supporting production workloads. The raw analysis indicates a *potential* monthly optimization opportunity of approximately **$4,606** through Compute Savings Plans, which offer significant discounts in exchange for commitment to consistent spend levels. This opportunity is available immediately with no infrastructure changes required, as the stable utilization pattern of production workloads makes them ideal candidates for commitment-based pricing.

A longer-term compute optimization opportunity exists through Graviton (ARM-based) instance migration, which could yield additional savings of approximately $1,800 per month. However, this requires architecture review and application compatibility testing before implementation and should be treated as a Phase 2 initiative rather than an immediate action.

6.3 Managed Database Configurations

Database services (RDS, MemoryDB, DocumentDB) represent approximately 28% of total spend, all running continuously with no Reserved Instance coverage. The raw analysis indicates a *potential* monthly optimization opportunity of approximately **$1,466** through Reserved Instance purchases: approximately $459 for RDS, $868 for MemoryDB, and $140 for DynamoDB. These opportunities are available immediately and represent low-risk, high-impact actions.

Deeper database optimization through instance rightsizing or cluster consolidation is possible but requires a comprehensive architecture review to ensure performance requirements are maintained. This should be treated as a follow-on opportunity after immediate pricing optimizations are implemented.

6.4 Storage Tiering and Lifecycle Management

Storage costs of approximately $156 per month include both S3 and EBS components. The raw analysis indicates a *potential* monthly optimization opportunity of approximately **$125** through two immediate actions: deletion of over 15 unused EBS volumes (approximately $75 per month) and migration from gp2 to gp3 volume types (approximately $50 per month). Both actions are low-effort, low-risk, and can be executed without architectural dependencies.

6.5 Discount and Commitment Coverage

The most significant opportunity identified is the implementation of commitment-based pricing across compute and database services. With 100% On-Demand spend and zero current coverage, the environment is positioned to benefit immediately from Savings Plans and Reserved Instances. The combined potential impact of commitment coverage across compute and databases is approximately **$6,072** per month, representing the highest-impact, lowest-risk opportunity available.

6.6 Interpreting Optimization Potential

In aggregate, the identified opportunities represent a potential monthly reduction of **$6,200–$8,200**, or approximately **25–32% of baseline spend** if fully realized and sustained. The conservative end of this range ($6,200) reflects immediate low-risk actions: Compute Savings Plans, Database Reserved Instances, and storage cleanup. The high end ($8,200) additionally includes Graviton migration, which is valid but requires a longer timeline and architectural groundwork.

These figures are directional indicators based on current-state patterns, not guaranteed outcomes. Actual results depend on implementation choices, workload requirements, and ongoing operational discipline. The percentage impact (25-32%) is often more meaningful than absolute dollars for understanding the significance of these opportunities relative to baseline spend.

---

7 Governance & FinOps Maturity Perspective

The current state of cost governance reflects an organization that has focused on building and operating a reliable travel technology platform but has not yet established formal FinOps practices. The absence of AWS Budgets, cost allocation tags, Reserved Instances, Savings Plans, or regular cost reviews indicates that cloud cost management has been entirely reactive—costs are paid as incurred without proactive optimization or forecasting.

This early maturity stage is not uncommon for organizations in growth phases, particularly those that have prioritized feature development and operational reliability over cost efficiency. The client's stated cost-sensitivity and desire to reduce spend now provides a natural catalyst for advancing governance capabilities. The good news is that the infrastructure is stable and well-run, meaning that cost optimization can proceed from a solid operational foundation rather than addressing underlying instability.

The recommended immediate governance action is the implementation of AWS Budgets at multiple levels: account level for overall spend visibility, service level to track compute and database costs specifically, and environment level to distinguish production from development and staging costs. This provides immediate financial visibility and alerting capability without requiring significant effort or architectural change. Combined with the pricing optimizations identified in the previous section, establishing AWS Budgets creates a foundation for sustained cost governance as the platform continues to grow.

The stability and predictability of current spend patterns provide a favorable foundation for formalizing cost governance practices. Costs are driven by legitimate business demand, not anomalies or waste, which means that optimization efforts can focus on pricing efficiency rather than usage reduction. This positions the organization to implement cost management as an operational capability rather than a one-time exercise.

---

8 Options & Next Steps

The findings of this AWS Cost Health Check support two broad approaches to ongoing cost management, each with distinct characteristics and resource requirements.

The first approach involves addressing the identified opportunities internally on a periodic basis. This would involve implementing Compute Savings Plans and Database Reserved Instances through internal engineering effort, cleaning up unused EBS volumes, and establishing AWS Budgets for ongoing visibility. This approach assumes that engineering leadership has capacity to prioritize these actions and that cross-functional alignment between technical and finance stakeholders can be achieved. The advantage is direct internal ownership; the consideration is that ongoing discipline requires sustained attention.

The second approach treats cost efficiency as a continuous operational capability with defined ownership and regular review cycles. This would involve establishing formal cost governance practices, implementing the pricing optimizations identified, and potentially engaging external support for deeper architectural optimization (such as the Graviton migration assessment). This approach is typically chosen when organizations prefer sustained oversight and reduced operational burden on internal teams, or when the scale of cloud spend justifies dedicated attention.

Both approaches are valid. The appropriate choice depends on organizational maturity, resource availability, and the desired balance between internal ownership and external support. The purpose of this report is to enable a clear and informed decision, not to prescribe a single outcome.
